

JUMMA MOHAMMAD TELI

Data Analyst | Power BI (DAX, Power Query, Service Admin) · Advanced SQL, Snowflake & Teradata · dbt & Dataform · Data Modelling · Data Quality · Advanced Python · Tableau, QuickSight, Looker, Omni & Qlik · Client & Wealth Analytics

Birmingham, UK | open to relocation & hybrid working | jummamohammad477@gmail.com | +44 7442 001088 | Right to work: Graduate visa (valid until Jan 2027)

linkedin.com/in/jumma-mohammad | github.com/jumma786 | jumma786.github.io/portfolio

PROFESSIONAL SUMMARY

Data analyst with **7+ years** across financial services, telecom and public-sector environments, currently at **UBS**, turning large and awkward datasets into reporting a business actually decides on. **Advanced Power BI** is my primary delivery surface — **DAX** (time intelligence, filter context, variables, iterators), **Power Query / M** with error handling that survives a refresh, **semantic models designed from scratch**, and full **Power BI Service administration**: workspaces, deployment pipelines, refresh schedules, gateways, row-level security and shared datasets. **Advanced SQL** underpins everything — 50+ tuned queries, views and stored procedures across **Snowflake**, an enterprise **Teradata** warehouse, Azure SQL, SQL Server and **BigQuery**, using CTEs and window functions over multi-million-row tables. I build the transformation layer in **dbt and Dataform** with **data quality tests inside the pipeline**, design **Kimball star schemas**, and treat data quality as a deliverable rather than a preamble — including an audit that recovered **296,834 records** a silent ETL fault had corrupted while every dashboard still looked healthy. I define **KPIs, calculation logic and data dictionaries** so a measure means the same thing everywhere it appears, extract delivery data from **Azure DevOps**, and automate recurring reporting with **Power Automate**, SQL and Python — saving ~8 hours/week. In financial services I **query Snowflake and Teradata to profile client segments** and run **ad-hoc deep dives into customer journeys**, reporting on **AUM/AUMA and client retention** and surfacing the **integration data anomalies** that would otherwise reach a client-facing number — then translate what the data says into **plain-English, commercially actionable recommendations** for non-technical business audiences. Also fluent in **Tableau, Amazon QuickSight, Looker, Omni and Qlik, Python**, and **Git-based delivery** with reviewed pull requests and CI. Backed by a public portfolio: a tested **dbt pipeline on 9.4M rows**, a **6.36M-record dimensional warehouse** with nine quality gates, live analytics products, and a **10-project MLOps portfolio with 132 passing tests**. **MSc Internet of Things and Future Networks with Merit** (Queen Mary University of London); **BSc Information Technology (First Class)**; IBM-certified in Data Analytics and Visualisation.

CORE SKILLS

Power BI: DAX, Power Query (M), Semantic Models, Dashboards
Advanced SQL: Complex Queries, CTEs, Window Functions, Tuning
Teradata: Teradata SQL, Query Tuning, Enterprise Data Warehouse
Data Transformation: dbt, Dataform — Models, Tests, Docs, Incremental
Data Quality: Validation, Profiling, Reconciliation, Automated Gates

Cloud Platforms: Snowflake, BigQuery, Databricks, AWS (S3, Redshift, QuickSight)
ETL/ELT & Orchestration: Apache Airflow, Talend, SSIS, dlt, Python
Automation: Power Automate, UiPath RPA, Scheduled Jobs, Python
Wealth & Client KPIs: AUM/AUMA, Client Retention, Attrition, Book Analysis
KPI & Metrics Frameworks: Definitions, Calculation Logic, Targets
Statistics & Experimentation: Regression, Forecasting, A/B Testing
AI-Assisted Development: Claude, Microsoft Copilot, Cursor, LLM Guardrails
Data Storytelling: Executive Narrative, Commercial Recommendations, Briefings

Power BI Service: Workspaces, Deployment Pipelines, Gateways, RLS
Snowflake: Snowflake SQL, Virtual Warehouses, Query Profiling, Cloud DWH
Databases: SQL Server, T-SQL, Azure SQL, PostgreSQL, MySQL, SSMS
Data Modelling: Kimball Star & Snowflake Schemas, Fact/Dimension Grain
Other BI Tools: Tableau, Amazon QuickSight, Looker (LookML), Omni, Qlik Sense, SSRS
Python: Pandas, NumPy, Scikit-learn, XGBoost, PySpark, pytest

Azure DevOps: Querying Work Item & Delivery Data, Boards, Backlogs
Version Control: Git, GitHub Actions CI/CD, Pull Requests, Peer Review
Customer Journey Analytics: Journey Deep Dives, Segment Profiling, Cohorts
Documentation & Governance: Data Dictionaries, Lineage, GDPR, FCA
Machine Learning & MLOps: Churn, SHAP, MLflow, Docker, FastAPI
Business Systems: Salesforce, GA4, Microsoft Fabric, Advanced Excel
Stakeholder Engagement: Requirements, UAT, Training, Agile/Scrum

WORK EXPERIENCE

Data Analyst (Contract, Remote) | UBS Oct 2024 – Present

- Design, build and maintain interactive **Power BI and Tableau dashboards** used daily by 20+ stakeholders and executives across risk, compliance and finance, built on **DAX measures, Power Query transformations and semantic models** designed for performance as data volumes grow.
- Administer the **Power BI Service** — workspaces and access, deployment pipelines promoting dev to test to production, refresh schedules and failure diagnosis, on-premises gateways, row-level security, dataflows and shared datasets.
- Write and tune **50+ complex SQL queries, views and stored procedures** across **Snowflake**, an enterprise **Teradata** warehouse, Azure SQL, SQL Server and BigQuery, using CTEs and window functions over multi-million-row tables to cut report execution times.
- Build and maintain **SQL models in dbt and Dataform**, layering staging, intermediate and mart models with **data quality tests inside the pipeline**, so business logic is defined once and a broken upstream assumption fails the build rather than reaching a dashboard.
- Query **Snowflake and Teradata environments to profile client segments** and run **ad-hoc deep dives into customer journeys**, reporting on **AUM/AUMA, client retention and attrition**, and identifying **integration data anomalies** where feeds disagree — before the discrepancy reaches a client-facing number.
- Extract, query and validate **delivery and operational data from Azure DevOps** alongside SQL sources, combining and transforming it into reliable datasets for repeatable reporting.
- Define **KPIs, calculation logic and reporting standards**, and maintain **data dictionaries, lineage and technical documentation** so measures stay consistent and another analyst can support what I build.
- Perform **data profiling, cleansing and validation** across millions of records, raising gaps, duplicates and figures that do not look right rather than publishing them as-is.
- Automate recurring reporting and extraction using **Power Automate**, SQL and Python, saving ~8 hours/week of manual effort while improving accuracy and consistency.
- Collaborate with risk, compliance, finance and product teams to translate business requirements into technical specifications, and turn complex findings into a clear narrative for non-technical business audiences — leading with the decision, then the number — so each analysis lands as a **strategic commercial recommendation** rather than a chart pack.
- Run **A/B and multivariate tests** and apply statistical methods — regression, forecasting, significance testing — fixing the success metric and decision rule before the result is known.
- Manage analytics under **Git version control** with branching, reviewed pull requests and CI, and use **Microsoft Copilot, Claude and Cursor** critically to accelerate development, documentation and analysis.

Data Analyst (Part-time) | IT Skills, London, UK Oct 2023 – Oct 2024

- Built and maintained **interactive Power BI dashboards** visualising KPIs across departments, increasing reporting efficiency by **30%** and enabling faster executive decision-making.
- Automated **ETL workflows in Talend, Python and SQL**, reducing manual processing time by **40%** and improving report delivery timeliness by **20%**.
- Integrated multiple online and offline sources into a **unified data warehouse**, developing shared, reusable datasets that replaced per-team extracts.
- Wrote and optimised complex SQL for **validation, profiling and cleansing**, holding **99%+ accuracy** across downstream reporting.
- Gathered requirements from department heads and defined metrics collaboratively, so reporting reflected agreed business logic rather than an analyst's assumption.
- Conducted **root-cause analysis on data discrepancies** and presented findings to senior management with actionable recommendations.
- Trained **50+ end users** onto reporting tools and maintained data quality documentation and standard operating procedures, driving self-serve adoption.

- Created **scheduled reports and Power BI Service subscriptions**, streamlining distribution across the organisation.

Data Analyst | Skyway Internet Services, Nagaur, India

Apr 2019 – Jul 2023

- Designed and operated SQL and Python **data pipelines processing 100K+ records daily**, accountable for reliability and timely availability to the reporting layer.
- Analysed large operational datasets using SQL and statistical techniques, surfacing **performance trends and anomalies** that drove measurable business improvements.
- Built **churn and forecasting models** for key commercial metrics, enabling evidence-based resource allocation and strategic planning.
- Developed **Power BI and Tableau dashboards** communicating KPIs to non-technical stakeholders, improving data literacy and reducing ad-hoc reporting requests by **35%**.
- Conducted **customer segmentation (RFM)** to identify high-value accounts and inform targeted retention campaigns, then measured what those campaigns returned.
- Ran **test-and-control campaign experiments**, so an observed lift could be attributed to the activity rather than to seasonality.
- Implemented **automated data validation and quality checks**, catching failures before they reached a dashboard rather than after someone queried the number.
- Partnered with product and operations teams to define KPIs** and reporting requirements, and **mentored junior analysts** on SQL best practice, report design and data-quality standards.

IT Technician | BSF (Border Security Force), Delhi, India

Apr 2016 – Jan 2019

- Operated and maintained **secure communication systems carrying sensitive operational data**, accountable for their confidentiality, integrity and availability.
- Maintained configuration and change logs and **documented technical procedures**, producing the audit trail that made system changes traceable and reviewable.
- Trained 30+ staff** on new systems and protocols, and managed inventory and equipment records, improving operational efficiency across the unit.

KEY PROJECTS

Analytics Engineering Pipeline on dbt & BigQuery — 9.4M Rows, 15 Tests · dbt · BigQuery · DuckDB · Airflow · GitHub Actions CI

A full **staging → intermediate → marts** dbt project over 9.4M records: star schema with a hashed surrogate key, an **incremental fact table**, a pre-aggregated BI mart and a Jinja macro for shared logic. **The 15 tests earned their place** — one month passed clean, the full dataset failed two, exposing 3 corrupt records and a duplicate key shipped in the source; both fixed at the boundary with **CI running a full dbt build on every push**. One codebase, two warehouses: DuckDB in dev, BigQuery in prod.

GitHub: [nyc_taxi_dbt](#)

Dimensional Warehouse with Nine Automated Quality Gates — 6.36M Records · SQL · Kimball Star Schema · Python ETL · PaySim Dataset

6.36M transactions from the PaySim simulator through a star schema to a BI-ready layer, with **9 automated data-quality gates** firing before anything reaches a consumer. Quantified an existing detection control's true recall at **0.19%**, then rebuilt it as a four-signal scorecard reaching **48.6% recall at 76.4% precision** — the analysis changed the control, not just the report. Documented afterwards as a formal requirements case study (BRD, user stories, traceability matrix).

GitHub: [enterprise-dwh-analytics · ba-requirements-fraud-scorecard-case-study](#)

Operational Analytics & the Data Quality Audit Behind It — 495,771 Records · Power BI · DAX · SQL Server Star Schema · Python ETL

On-time performance, delay-cause and cancellation reporting on **495,771 records** over a SQL Server star schema with a Power BI executive dashboard. **The audit is the point**: the original import ran without errors and the dashboards looked healthy, while **296,834 rows (60%) carried a null date** from a silent format fault and **9,056 cancellations were counted as on-time**. Recovering them changed the reported numbers.

GitHub: [airline-operations-analytics](#)

Live Performance Dashboard Against a National Target — 41 Months, 42 Regions · Python · Forecasting · Geospatial · Streamlit

A deployed analytics product tracking **performance against a published national standard**: trends, specialty breakdowns, exception tracking on long waiters, forecasting and geographic inequality maps. Reports the honest position — **65.3% within standard against a 92% target** — and surfaces a **21-point gap** between best and worst region, with its data quality caveats documented on the face of it.

GitHub: [NHS-RTT-Dashboard · live app linked in the README](#)

Retail Analytics & Churn — 1.07M Transactions · SQL Server · RFM Segmentation · 4-Page Power BI · XGBoost

End-to-end analysis of **1.07M UK retail transactions**: full cleaning, T-SQL analysis, **RFM behavioural segmentation**, a **4-page Power BI dashboard**, and a churn model on leakage-safe time-based features identifying **£1.26M of revenue exposure** across the top 100 at-risk customers.

GitHub: [online-retail-analytics](#)

Attribution Modelling on 16.5M Events · Python · Markov Chains · Shapley · dbt · LookML · 30 Tests

Multi-touch attribution across **16,468,027 impressions, 6.1M users and 675 campaigns** on real Criteo data: Markov removal-effect and Shapley value, reallocating **£468,342 of a £5M budget** — with the honest caveat that only **39.7% of conversions are multi-touch**. Engine fully run and reproducible (30 passing tests); the dbt marts and LookML layer are written but not yet run against a live warehouse.

GitHub: [marketing-attribution-modelling](#)

Hospital Analytics End-to-End — 27,891 Encounters · T-SQL · 3-Page Power BI · DAX · XGBoost · SHAP · FastAPI

Schema design, data-quality checks and T-SQL analytics on SQL Server, a **3-page Power BI report** (Operations, Financial, Patient Risk) built on DAX measures, and an XGBoost readmission model reaching **AUC 0.907** with SHAP explainability, deployed as a live FastAPI service. Dataset is **Synthea synthetic EHR**, stated plainly in the repository.

GitHub: [hospital-analytics-end-to-end · hospital-readmission-api](#)

Self-Serve Analytics Agent — Plain English to Governed SQL · Python · LangGraph · FastAPI · PostgreSQL · 133 Tests

A multi-agent pipeline letting non-technical users query a database in plain English — converting the question to SQL, **refusing to execute anything that is not read-only**, running it and explaining the result, with JWT auth, RAG document Q&A and PDF reporting. The engineering constrains what the model may touch.

GitHub: [ai-data-analysis-agent](#)

Power BI & Tableau Dashboard Portfolio · Power BI (DAX) · Tableau · Dynamic KPIs · Drill-Through

Published dashboards across operational domains: **supply chain** cost, supplier performance and route analysis; **dispatch delays and product returns** for a UK e-commerce operator; **workforce attrition**; UK retail revenue, profit and segments; and Tableau work on housing and salary distribution. One discipline throughout — establish what the reader must decide, then design backwards from that decision.

GitHub: [supply-chain-performance-powerbi · hr-attrition-analysis-powerbi · house-price-analysis-tableau](#)

Early-Warning Risk Scorecard — 824,364 Transactions · Python · Exposure Metrics · Transparent Scoring · HTML Dashboard

Turns real transaction data into **exposure metrics, dilution and concentration signals** and a **transparent early-warning scorecard**, served through a self-contained dashboard, with **every score tracing back to the inputs that produced it** — an opaque risk number gets ignored by the people who must act on it.

GitHub: [invoice-finance-risk-monitor](#)

End-to-End MLOps Portfolio — 10 Projects, 132 Passing Tests · MLflow · FastAPI · Docker · Airflow · Kubernetes · DVC · Optuna

Ten production-style projects across the full ML lifecycle: model tournament, scheduled retraining (DVC), tuning (Optuna), FastAPI + Docker serving, feature store, **drift detection (PSI/KS-test)**, A/B testing, Airflow orchestration and a Kubernetes platform — **132 passing tests** on real UCI datasets.

GitHub: [github.com/jumma786 \(mlops-model-tournament, mlops-model-serving, mlops-k8s-platform + 7 more\)](#)

EDUCATION

MSc Internet of Things and Future Networks — Merit

Queen Mary University of London, UK · Sep 2023 – Sep 2024

BSc Information Technology — First Class

Sabarmati University, Gujarat, India · Sep 2012 – Jun 2015

CERTIFICATIONS

IBM: Data Analytics, Tools for Data Science, Data Visualization & Dashboards (Excel, Cognos), Introduction to AI · **HackerRank**: SQL · **LinkedIn**: Power BI for Finance, Data Science Methodologies